

Claude is a space to think

There are many good places for advertising. A conversation with Claude is not one of them.

Advertising drives competition, helps people discover new products, and allows services like email and social media to be offered for free. We've run our own [ad campaigns](#), and our AI models have, in turn, helped many of our customers in the advertising industry.

But including ads in conversations with Claude would be incompatible with what we want Claude to be: a genuinely helpful assistant for work and for deep thinking.

We want Claude to act unambiguously in our users' interests. So we've made a choice: Claude will remain ad-free. Our users won't see "sponsored" links adjacent to their conversations with Claude; nor will Claude's responses be influenced by advertisers or include third-party product placements our users did not ask for.

The nature of AI conversations

When people use search engines or social media, they've come to expect a mixture of organic and sponsored content. Filtering signal from noise is part of the interaction.

Conversations with AI assistants are meaningfully different. The format is open-ended; users often share context and reveal more than they would in a search query. This openness is part of what makes conversations with AI valuable, but it's also what makes them susceptible to influence in ways that other digital products are not.

Our [analysis of conversations](#) with Claude (conducted in a way that keeps

all data [private and anonymous](#)) shows that an appreciable portion involve topics that are sensitive or deeply personal—the kinds of conversations you might have with a trusted advisor. Many other uses involve complex software engineering tasks, deep work, or thinking through difficult problems. The appearance of ads in these contexts would feel incongruous—and, in many cases, inappropriate.

We still have much to learn about the impact of AI models on the people who use them. [Early research](#) suggests both benefits—like people finding support they couldn't access elsewhere—and risks, including the potential for models to reinforce harmful beliefs in vulnerable users. Introducing advertising incentives at this stage would add another level of complexity. [Our understanding](#) of how models translate the goals we set them into specific behaviors is still developing; an ad-based system could therefore have unpredictable results.

Incentive structures

Being genuinely helpful is one of the core principles of [Claude's Constitution](#), the document that describes our vision for Claude's character and guides how we train the model. An advertising-based business model would introduce incentives that could work against this principle.

Consider a concrete example. A user mentions they're having trouble sleeping. An assistant without advertising incentives would explore the various potential causes—stress, environment, habits, and so on—based on what might be most insightful to the user. An ad-supported assistant has an additional consideration: whether the conversation presents an opportunity to make a transaction. These objectives may often align—but not always. And, unlike a list of search results, ads that influence a model's responses may make it difficult to tell whether a given recommendation comes with a commercial motive or not. Users shouldn't have to second-

guess whether an AI is genuinely helping them or subtly steering the conversation towards something monetizable.

Even ads that don't directly influence an AI model's responses and instead appear separately within the chat window would compromise what we want Claude to be: a clear space to think and work. Such ads would also introduce an incentive to optimize for engagement—for the amount of time people spend using Claude and how often they return. These metrics aren't necessarily aligned with being genuinely helpful. The most useful AI interaction might be a short one, or one that resolves the user's request without prompting further conversation.

We recognize that not all advertising implementations are equivalent. More transparent or opt-in approaches—where users explicitly choose to see sponsored content—might avoid some of the concerns outlined above. But the history of ad-supported products suggests that advertising incentives, once introduced, tend to expand over time as they become integrated into revenue targets and product development, blurring boundaries that were once more clear-cut. We've chosen not to introduce these dynamics into Claude.

Our approach

Anthropic is focused on businesses, developers, and helping our users flourish. Our business model is straightforward: we generate revenue through enterprise contracts and paid subscriptions, and we reinvest that revenue into improving Claude for our users. This is a choice with tradeoffs, and we respect that other AI companies might reasonably reach different conclusions.

Expanding access to Claude is central to our public benefit mission, and we want to do it without selling our users' attention or data to advertisers. To that end, we've [brought AI tools and training to educators](#) in over 60 countries, begun national AI education pilots with [multiple governments](#),

and made Claude [available to nonprofits](#) at a significant discount. We continue to invest in our smaller models so that our free offering remains at the frontier of intelligence, and we may consider lower-cost subscription tiers and regional pricing where there is clear demand for it. Should we need to revisit this approach, we'll be transparent about our reasons for doing so.

Supporting commerce

AI will increasingly interact with commerce, and we look forward to supporting this in ways that help our users. We're particularly interested in the potential of agentic commerce, where Claude acts on a user's behalf to handle a purchase or booking end to end. And we'll continue to build features that enable our users to find, compare, or buy products, connect with businesses, and more—when they choose to do so.

We're also exploring more ways to make Claude a focused space to be at your most productive. Users can already [connect third-party tools](#) they use for work—like Figma, Asana, and Canva—and interact with them directly within Claude. We expect to introduce many more useful integrations and expand this toolkit over time.

All third-party interactions will be grounded in the same overarching design principle: they should be initiated by the *user* (where the AI is working for them) rather than an *advertiser* (where the AI is working, at least in part, for someone else). Today, whether someone asks Claude to research running shoes, compare mortgage rates, or recommend a restaurant for a special occasion, Claude's only incentive is to give a helpful answer. We'd like to preserve that.

A trusted tool for thought

We want our users to trust Claude to help them keep thinking—about their work, their challenges, and their ideas.

Our experience of using the internet has made it easy to assume that advertising on the products we use is inevitable. But open a notebook, pick up a well-crafted tool, or stand in front of a clean chalkboard, and there are no ads in sight.

We think Claude should work the same way.